

# Introduction to Data & Numerical Summaries (continued)

Emile Latour, Nicky Wakim, Meike Niederhausen

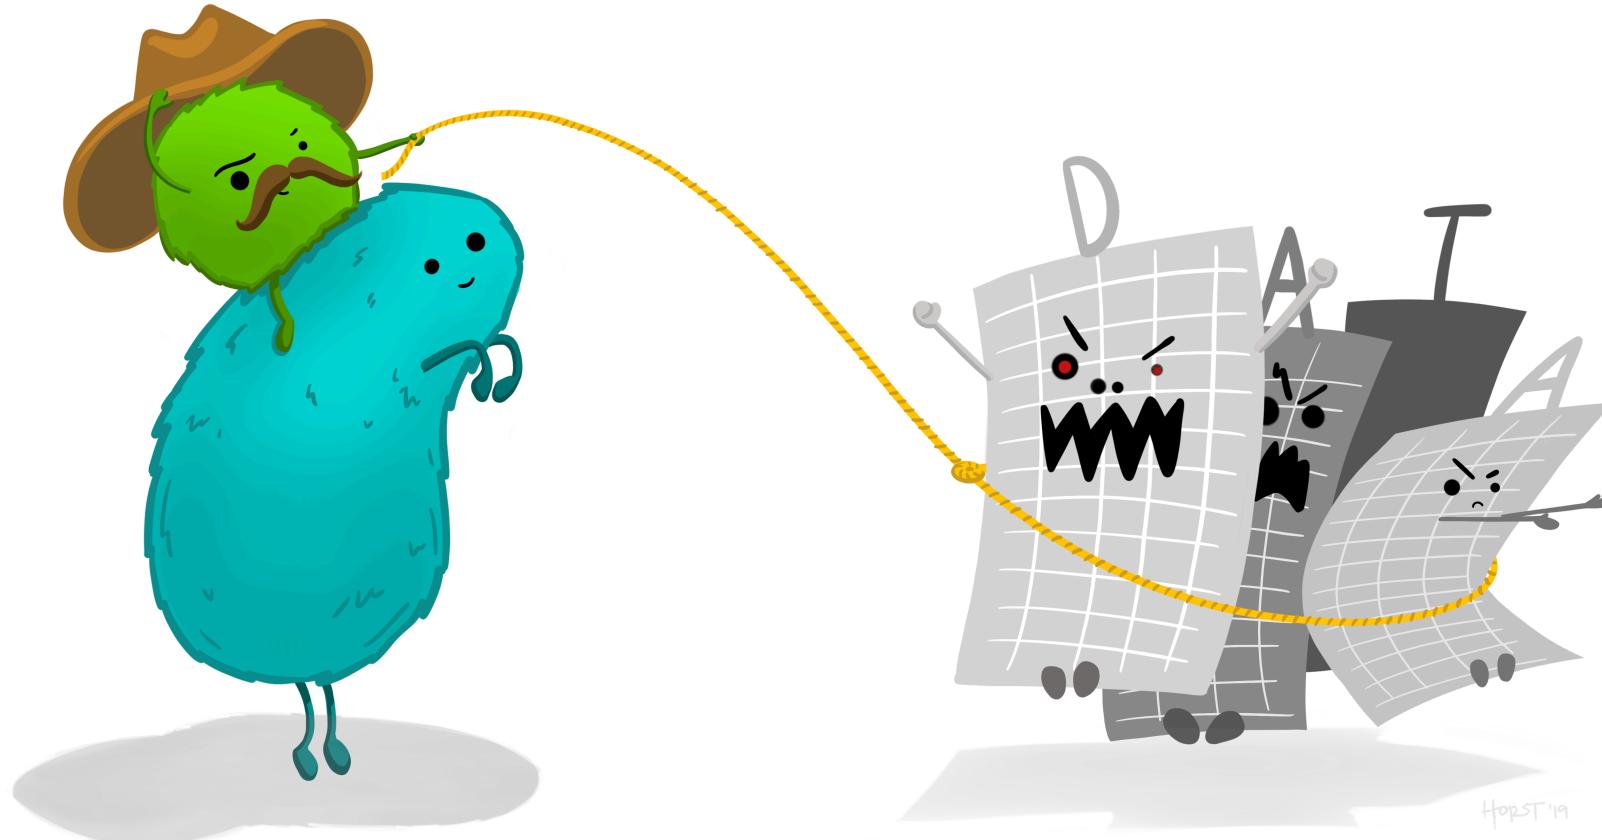
January 5, 2026

# Learning objectives (today)

By the end of class, you should be able to:

1. Identify observations and variables in a dataset
2. Distinguish numerical and categorical variables (discrete/continuous; nominal/ordinal)
3. Describe the center of numerical data using the mean and median
4. Describe the spread of numerical data using the SD and IQR
5. Explain why skewness and outliers affect some summaries more than others

## Intro to Data (1.2)



Artwork by @allison\_horst

## A motivating example: frog reproduction data<sup>1</sup>

To make ideas about data more concrete, we will work with a dataset from an evolutionary biology study on frog reproduction.

Female frogs invest energy into reproduction by producing eggs. Because energy is limited, there are natural trade-offs: - producing **many small eggs** versus - producing **fewer larger eggs**

Researchers collected data from frog egg clutches found at breeding ponds across multiple locations that differed in altitude.

The goal was to understand how reproductive investment varies across environments.

1. Dataset adapted from an example in *Introductory Statistics for the Life and Biomedical Sciences*.

## The frog study data (1/2)

<b>row</b>	<b>altitude</b>	<b>egg.size</b>	<b>clutch.size</b>	<b>clutch.volume</b>	<b>body.size</b>
1	3,462.00	1.95	181.97	177.83	3.63
2	3,462.00	1.95	269.15	257.04	3.63
3	3,462.00	1.95	158.49	151.36	3.72
150	2,597.00	2.24	537.03	776.25	NA

## The frog study data (2/2)

row	altitude	egg.size	clutch.size	clutch.volume	body.size
1	3,462.00	1.95	181.97	177.83	3.63
2	3,462.00	1.95	269.15	257.04	3.63
3	3,462.00	1.95	158.49	151.36	3.72
150	2,597.00	2.24	537.03	776.25	NA

- Each **observation** is a **row**.
- Each **variable** is a **column**.
- Each **cell** contains a **single value**.
- All the **observations** and **variables** together make a **data frame**. The textbook will sometimes call it a data matrix.
  
- **Missing values:** NA here means “missing”. The clutch #150 is missing `body.size`.

## The frog study variables and descriptions

variable	description
altitude	Altitude of the study site (meters above sea level)
latitude	Latitude of the study site (degrees)
egg.size	Average diameter of an individual egg (mm)
clutch.size	Estimated number of eggs in a clutch
clutch.volume	Total volume of the egg clutch ( $\text{mm}^3$ )
body.size	Body length of the egg-laying female frog (cm)

## Types of variables

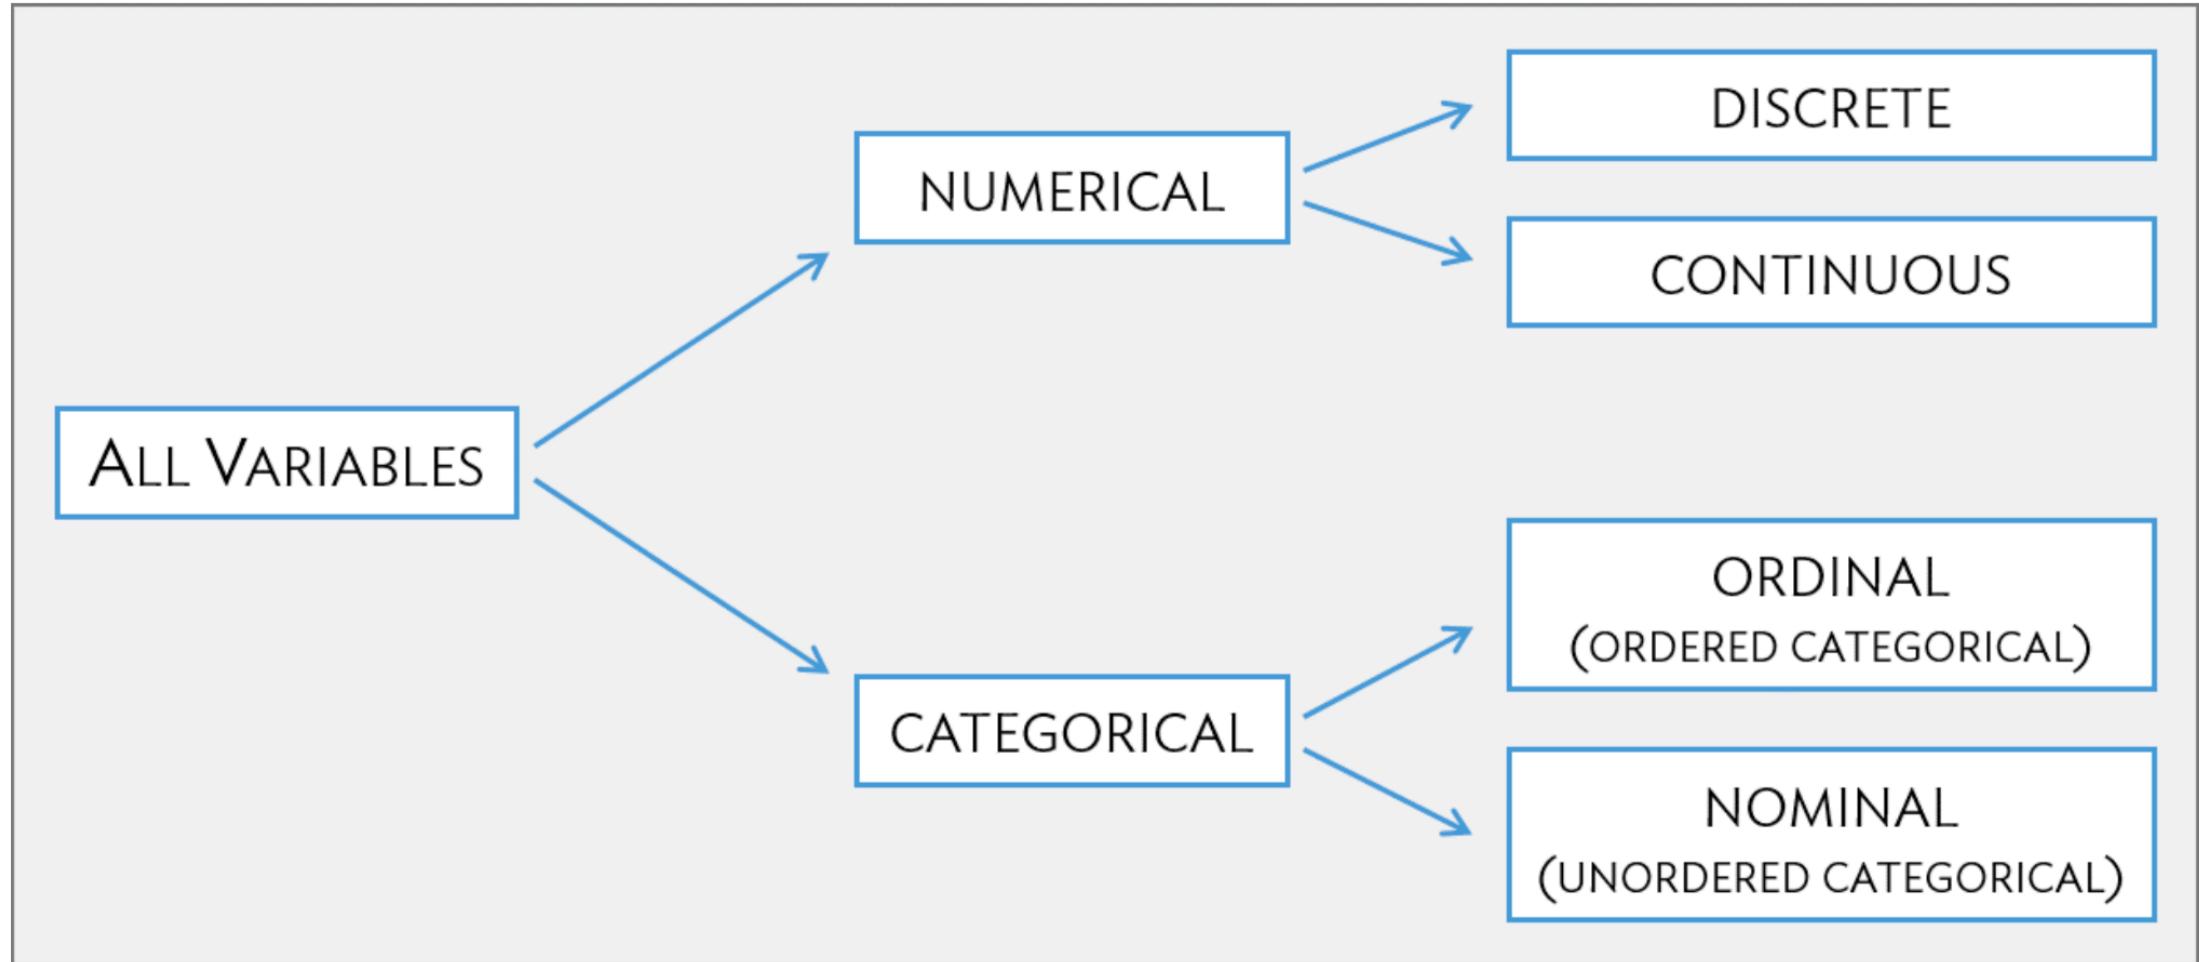


Figure 1.8: Breakdown of variables into their respective types.

# Numerical variables

**Numerical variables** take on values for which mathematical operations (addition, subtraction, averaging) are meaningful.

There are two common types:

## Discrete

- Take on **separate, countable values**
- Usually whole numbers
- Arise from counting

### Examples:

- Number of eggs in a clutch
- Number of clinic visits
- Number of children in a family

## Continuous

- Can take on **any value within a range**
- Often measured rather than counted
- Decimals are possible

### Examples:

- Egg size (mm)
- Height or weight
- Blood pressure

# Categorical variables

**Categorical variables** describe qualities or group membership rather than numerical magnitude. The values are called **levels**.

## Nominal

- Categories have **no natural ordering**
- Differences are labels, not amounts

### Examples:

- Species
- Blood type
- Gender identity

## Ordinal

- Categories have a **meaningful order**
- The spacing between categories is not necessarily equal

### Examples:

- Disease severity (mild, moderate, severe)
- Likert-scale responses (strongly disagree -> strongly agree)
- Education level

# How are data stored, how do we use them?

- Often, data are in an Excel sheet, or a plain text file (.csv, .txt)
- .csv files open in Excel automatically, but actually are plain text
- Usually, columns are variables/measures and rows are observations (i.e. a person's measurements)

## Data in R

- We can import data from many file types, including .csv, .txt., and .xlsx
  - We will cover this on a later date
- Once imported, R typically stores data as **data frames**, or **tibbles** if using the **tidyverse** package (more on this later).
  - For our purposes, these are essentially the same, and I will tend to use the terms interchangeably.
  - These are examples of what we call **object types** in R.

## Which type is each frog variable?

variable	description	type
altitude	Altitude of the study site (meters above sea level)	Numerical (continuous)
latitude	Latitude of the study site (degrees)	Numerical (continuous)
egg.size	Average diameter of an individual egg (mm)	Numerical (continuous)
clutch.size	Estimated number of eggs in a clutch	Numerical (discrete)
clutch.volume	Total volume of the egg clutch (mm <sup>3</sup> )	Numerical (continuous)
body.size	Body length of the egg-laying female frog (cm)	Numerical (continuous)

# How variable types appear in R (preview)

We are **not** using R in depth yet.

This slide is here to help you later connect: - the variable types we talk about conceptually, and - how R stores those variables internally.

You do **not** need to memorize this now.

R type	Variable type	What it represents
integer	Numerical (discrete)	Whole-number counts
double / numeric	Numerical (continuous)	Measured values with decimals
factor	Categorical	Categories stored with defined levels
character	Categorical	Text labels
logical	Categorical	TRUE / FALSE values

## Summarizing numerical data (1.4)

Once data have been collected, a first step is to **describe what we see**.

For **numerical variables**, we often want to understand:

- what values are *typical*, and
- how much the values *vary* across observations.

Numerical summaries help us:

- quickly describe large datasets,
- compare groups,
- and communicate patterns clearly.

In this section, we focus on summaries for **numerical variables**. We will return to categorical variables later.

# Two questions we often ask

For a numerical variable, two questions come up repeatedly:

## 1. Where are the values centered?

- What is a “typical” value?

## 2. How spread out are the values?

- Are observations tightly clustered or widely dispersed?

The summaries we introduce next are designed to answer these two questions.

### Note

Some examples will show R output as a reference. You do not need to understand the R code yet — focus on the ideas.

## Measures of center: the mean

The **mean** (or average) describes the center of a numerical variable.

It is calculated by:

- adding all observed values, and
- dividing by the number of observations.

$$\bar{x} = \frac{x_1 + x_2 + \cdots + x_n}{n} = \sum_{i=1}^n \frac{x_i}{n}$$

where:

- $x_1, x_2, \dots, x_n$  are the observed values, and
- $n$  is the number of observations.

The mean uses **all** observed values and is sensitive to unusually large or small values. All data values contribute equally, even “outliers.”

## Example: mean clutch volume

Using the frog data, we can ask:

**What is a typical clutch volume?**

For this dataset:

- each observation is one egg clutch
- clutch volume is measured in mm<sup>3</sup>

$$\bar{x} = \frac{177.8 + 257.0 + \dots + 933.3}{431} = \frac{380346.3}{431} \approx 882.5$$

```
1 mean(frog$clutch.volume, na.rm = TRUE)
```

```
[1] 882.474
```

The mean clutch volume is 882.5 mm<sup>3</sup>.

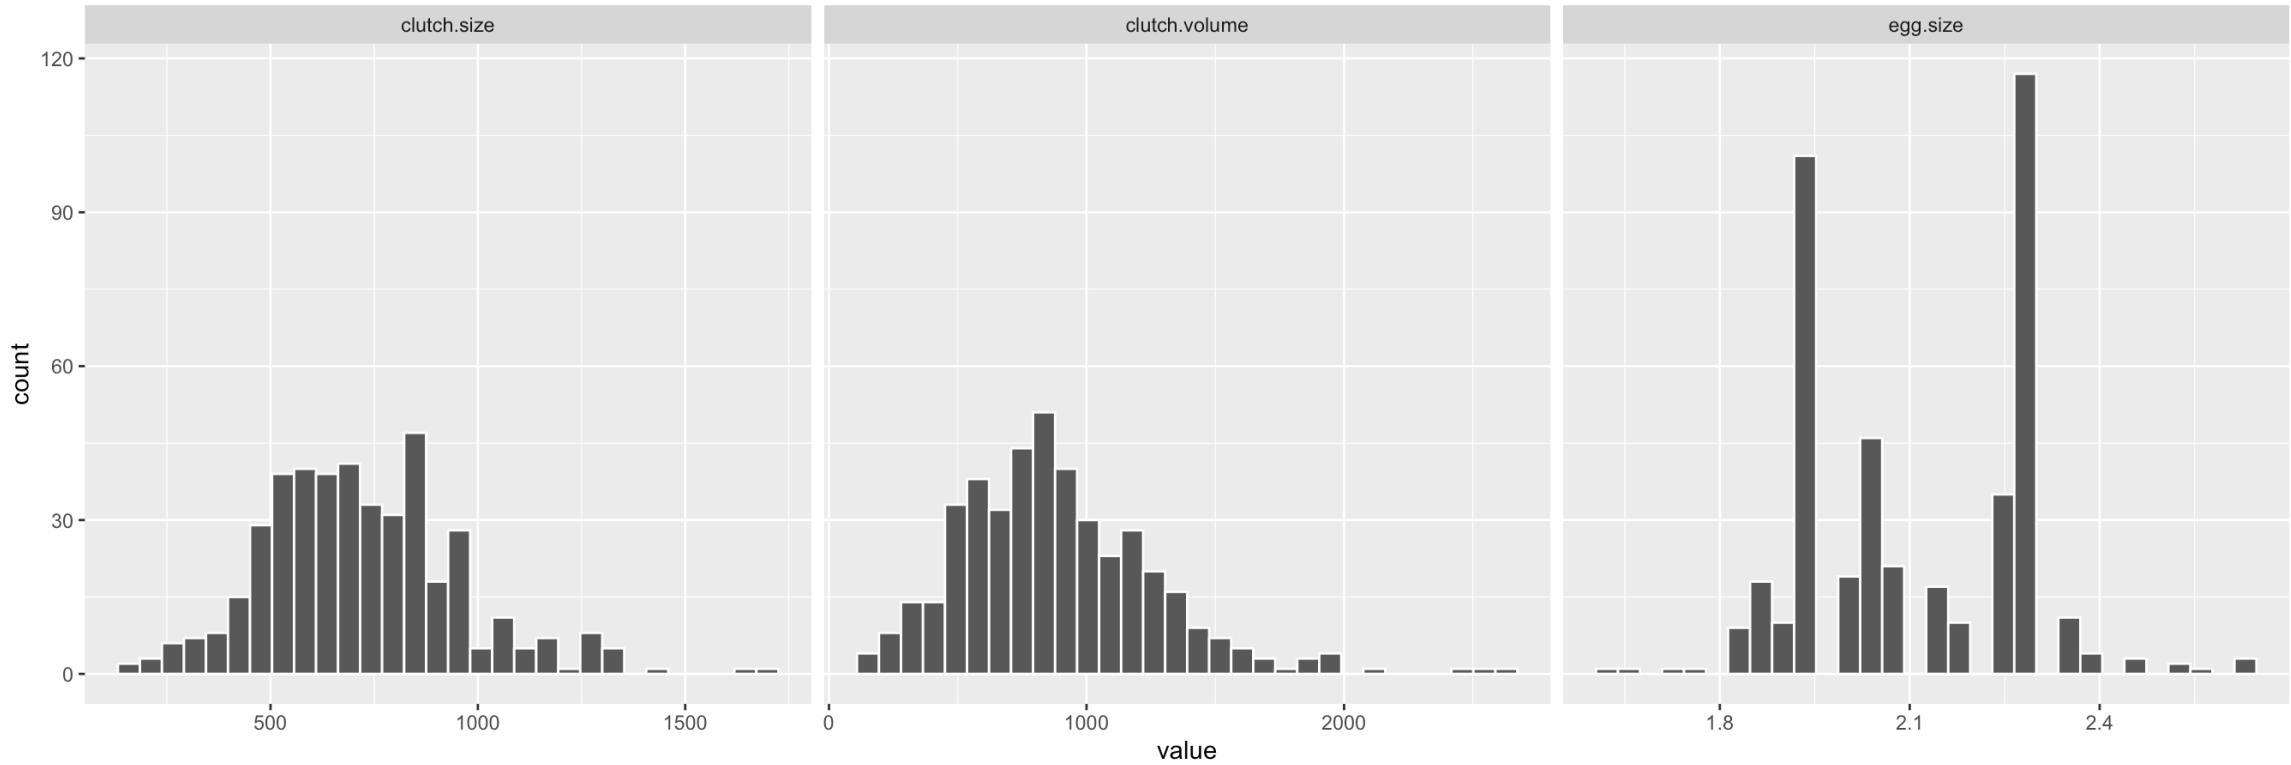
The mean tells us where the data balance, but it doesn't tell us how spread out they are.

## Things to note

- The mean of a binary variable with 0/1 coding gives the proportion (relative frequency) of ones in the sample.
- The mean can be heavily influenced by “outliers.”
  - $2, 7, 9 \rightarrow \bar{x} = \frac{2+7+9}{3} = 6$
  - $2, 7, 29 \rightarrow \bar{x} = \frac{2+7+29}{3} \approx 12.7$

# Distributions of frog variables

Histograms of three variables in frog dataset



## Measures of center: the median

The **median** describes the *center* of a numerical variable based on **order**, not magnitude.

- The median is the **middle value** once observations are ordered from smallest to largest.
- Half of the observations lie **below** the median
- Half of the observations lie **above** the median

If there are:

- an **odd** number of observations, the median is the middle value
- an **even** number of observations, the median is the average of the two middle values

Unlike the mean, the median is **not affected** by extreme values.

## Example: mean vs. median

Consider these two datasets:

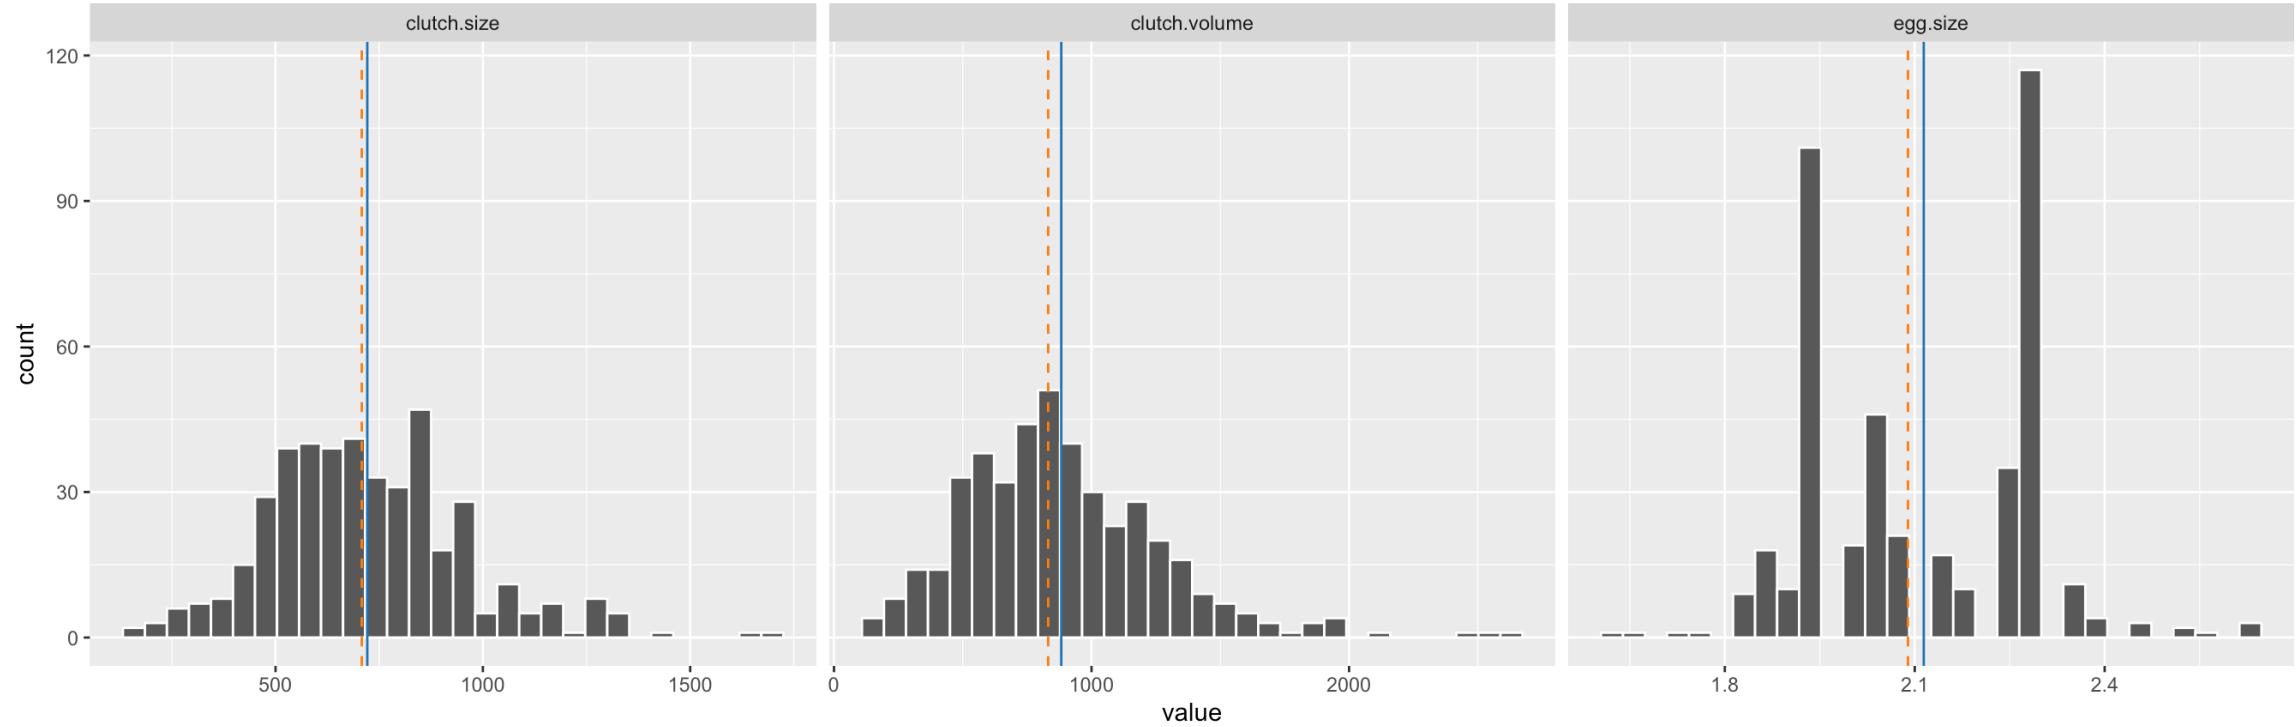
- 2, 7, 9
  - Mean = 6
  - Median = 7
- 2, 7, 29
  - Mean = 12.7
  - Median = 7

Only one value changed, but the **mean** changed a lot. The **median** did not.

# Distributions of frog variables

Histograms of three variables in frog dataset

Mean (solid blue), Median (dashed orange)



- When distributions are skewed, the **mean** is pulled toward the tail.
- The **median** stays closer to where most observations lie.

## A note on the mode

The **mode** is the most frequent value in a dataset.

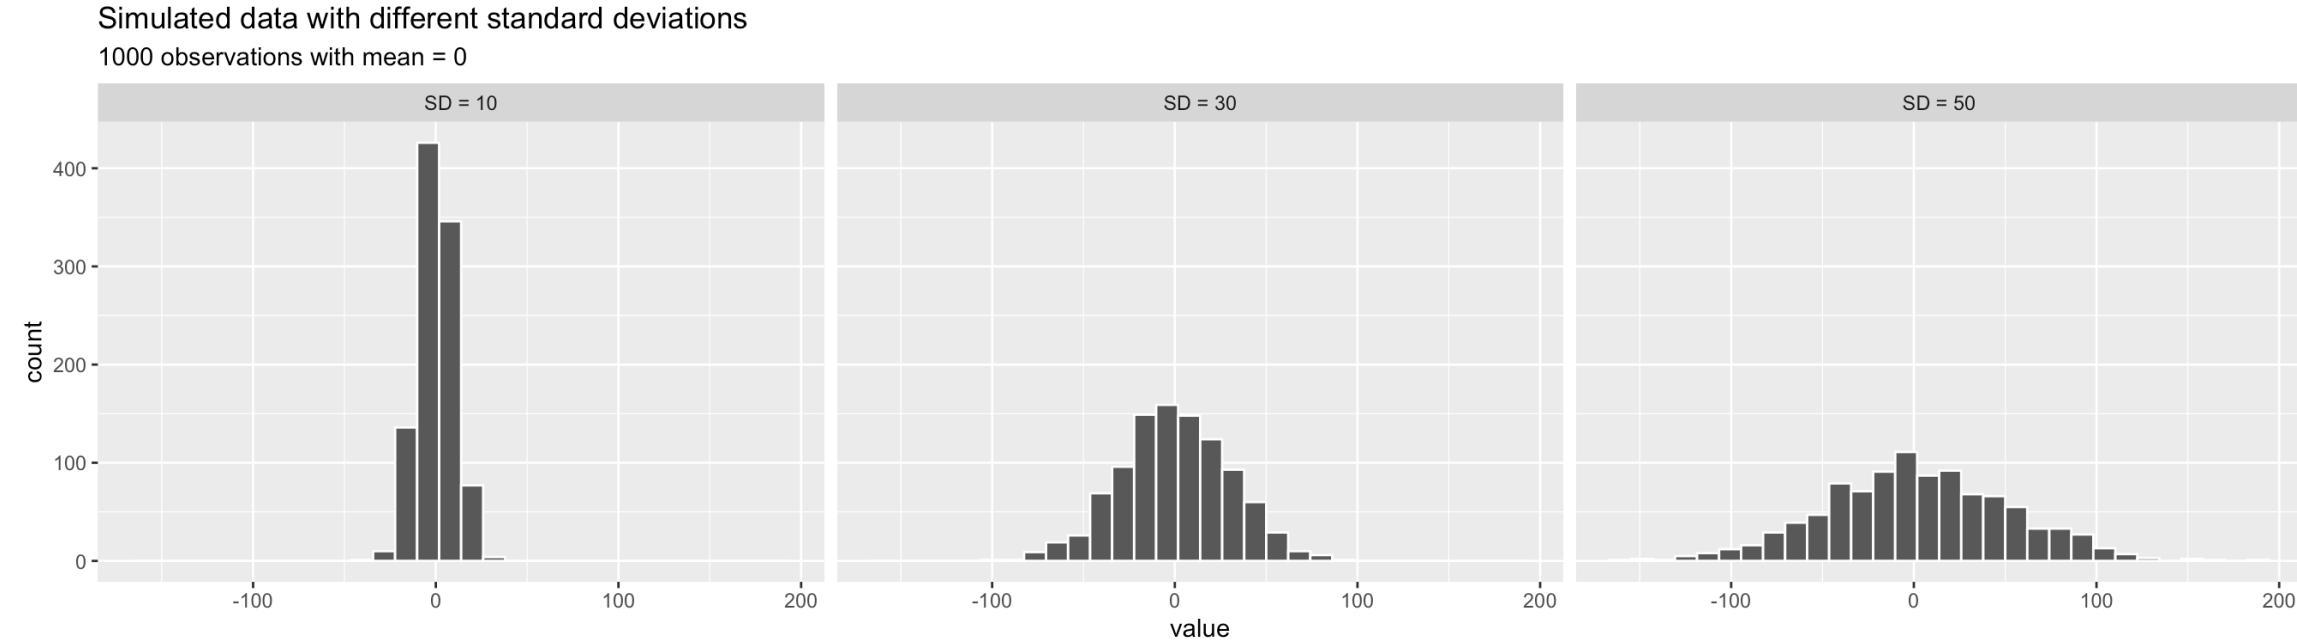
- For **numerical** data, the mode is often unstable or uninformative
- It depends on how values are measured or grouped
- As a result, it is rarely reported for continuous variables

The mode is more useful for **categorical** data, which we will revisit later.

## Measures of spread: standard deviation (SD) (1/3)

standard deviation is (approximately) the average distance between an observation and the mean

- An observation's **deviation** is the distance between its value  $x$  and the sample mean  $\bar{x}$ : deviation =  $x - \bar{x}$ .



## Measures of spread: SD (2/3)

- The **sample variance**  $s^2$  is the sum of squared deviations divided by the number of observations minus 1.

$$s^2 = \frac{(x_1 - \bar{x})^2 + (x_2 - \bar{x})^2 + \cdots + (x_n - \bar{x})^2}{n - 1} = \sum_{i=1}^n \frac{(x_i - \bar{x})^2}{n - 1}$$

where  $x_1, x_2, \dots, x_n$  represent the  $n$  observed values.

- The **standard deviation**  $s$  (or  $sd$ ) is the square root of the variance.

$$s = \sqrt{s^2} = \sqrt{\frac{1}{n - 1} \sum_{i=1}^n (x_i - \bar{x})^2}$$

## Measures of spread: SD (3/3)

Let's calculate the sample standard deviation for the `clutch.volume` from the frog dataset.

- Doing this by hand can be really time consuming!
- R does this easily for us.

```
1 sd(frog$clutch.volume, na.rm = TRUE)  
[1] 379.0527
```

So the standard deviation is 379.1 mm<sup>3</sup>.

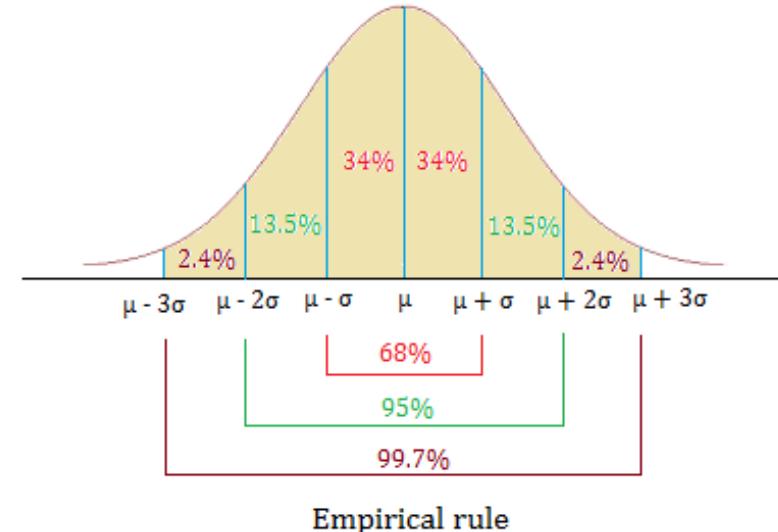
For the sample of 431 frog clutches, the mean clutch volume was 882.5 mm<sup>3</sup> (SD = 379.1 mm<sup>3</sup>).

# Empirical Rule: one way to think about the SD

For symmetric bell-shaped data, about

- 68% of the data are within 1 SD of the mean
- 95% of the data are within 2 SD's of the mean
- 99.7% of the data are within 3 SD's of the mean

These percentages are based off of percentages of a true normal distribution.



<https://statistics-made-easy.com/empirical-rule/>

## Measures of spread: interquartile range (IQR) (1/2)

The  $p^{th}$  percentile is the observation such that  $p\%$  of the remaining observations fall below this observation.

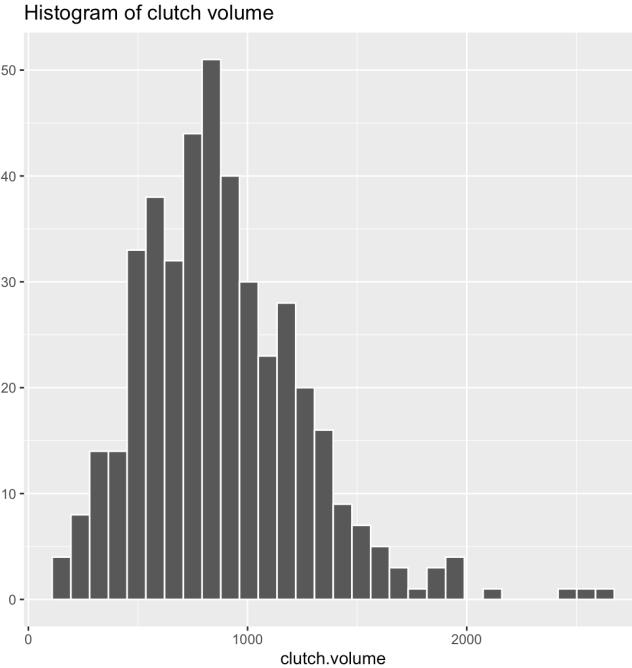
- The *first quartile*  $Q_1$  is the  $25^{th}$  percentile.
- The *second quartile*  $Q_2$ , i.e., the median, is the  $50^{th}$  percentile.
- The *third quartile*  $Q_3$  is the  $75^{th}$  percentile.

The **interquartile range (IQR)** is the distance between the third and first quartiles.

$$IQR = Q_3 - Q_1$$

- IQR focuses on the **middle 50% of the data**.
- Robust against outliers.

## Measures of spread: interquartile range (IQR) (2/2)



```
1 summary(frog$clutch.volume)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	151.4	609.6	831.8	882.5	1096.5	2630.3

$$IQR = Q_3 - Q_1 = 1096.5 - 609.6 = 486.9$$

```
1 IQR(frog$clutch.volume)
```

```
[1] 486.9009
```

## Measures of spread: range

The **range** describes spread using only the **smallest** and **largest** values.

$$\text{Range} = \max(x) - \min(x)$$

- Easy to compute and easy to interpret
- **Not robust:** it depends only on two observations
- A single unusually large or small value can dramatically change the range

### Note

Because it is so sensitive to extreme values, the range is usually less informative than the SD or IQR.

## Range and outliers (example)

- 2, 7, 9
  - $\min = 2$ ,  $\max = 9$
  - $\text{Range} = 9 - 2 = 7$
- 2, 7, 29
  - $\min = 2$ ,  $\max = 29$
  - $\text{Range} = 29 - 2 = 27$

Only one value changed, but the **range** changed a lot.

## Robust estimates

Summary statistics are called **robust estimates** if extreme observations (outliers) have little effect on their values

Estimate	Robust?
Sample mean	✗
Median	✓
Standard deviation	✗
IQR	✓
Range	✗

- For samples with extreme values or skewed distributions, the **median and IQR** often provide a more stable summary of center and spread than the mean, standard deviation, or range.