

R Module 2: Factors and Reshaping Data

BMSC 620 | R Workflow Series

Emile Latour



What We're Learning Today

By the end of this module, you will be able to:

1. Understand the difference between **characters** and **factors** in R
2. Use **forcats** functions to reorder and relabel factor levels
3. Recognize when data is in **wide** vs. **long** format
4. Use `pivot_longer()` and `pivot_wider()` to reshape data

Note

These skills come up constantly — every time you make a plot with categories or need to restructure data for analysis.



Two Problems You've Probably Already Hit

Problem 1: You make a bar chart and the categories are in alphabetical order instead of a meaningful order.

Problem 2: You have data with multiple columns that should really be rows (or vice versa) and your `ggplot()` code won't cooperate.

Both problems have the same root cause: R doesn't know what you know about the **structure** of your data.

Today we fix that.

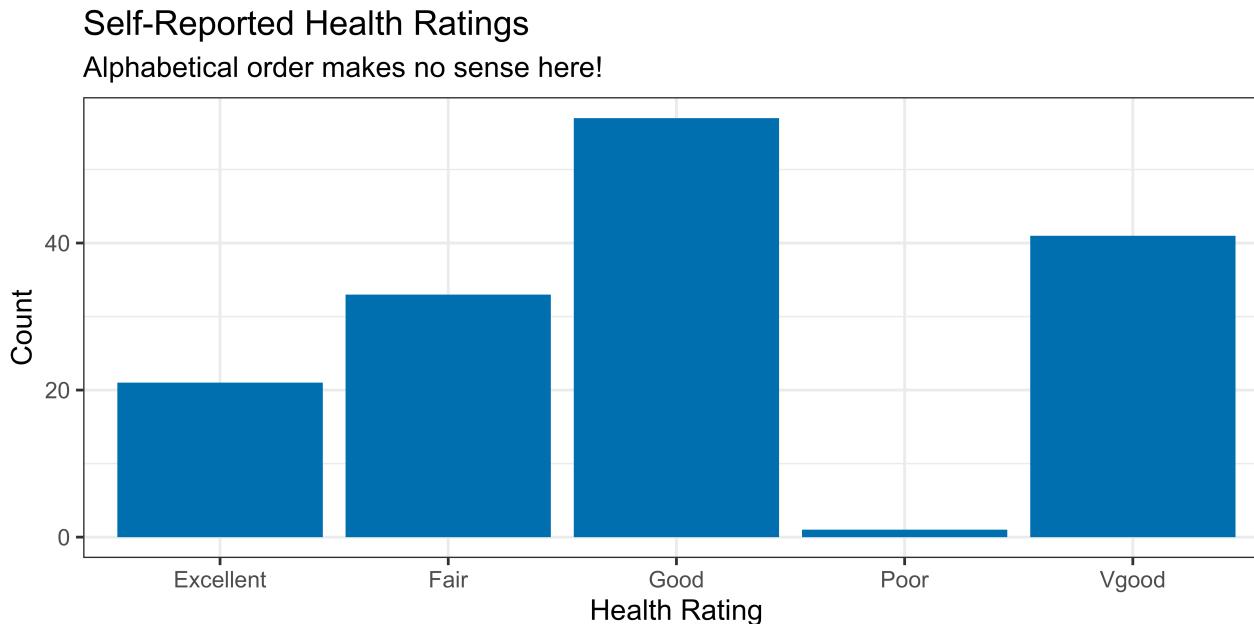


Part 1: What Is a Factor?



The Motivating Bug

```
1 nhanes.samp %>%
2   filter(!is.na(HealthGen)) %>%
3   ggplot(aes(x = HealthGen)) +
4   geom_bar(fill = "#0072B2") +
5   labs(title = "Self-Reported Health Ratings",
6        subtitle = "Alphabetical order makes no sense here!",
7        x = "Health Rating",
8        y = "Count") +
9   theme_bw(base_size = 16)
```



The categories are in **alphabetical order**. That's not what we want.



Character vs. Factor

Character

- Just text — R treats it like any string
- No inherent order
- Plotted alphabetically by default

```
1 class(nhanes.samp$HealthGen)
[1] "character"
```

Factor

- A categorical variable with a **defined set of levels**
- Levels have an **order** (even if it's arbitrary)
- You control the order

```
1 health_factor <- factor(nhanes.samp$HealthG
2 class(health_factor)
[1] "factor"
1 levels(health_factor)
[1] "Excellent" "Fair"      "Good"       "Poor"
"Vgood"
```

When R converts a character to a factor automatically, it uses **alphabetical order** — which is almost never what you want.



Creating Factors with `factor()`

You can specify the exact order of levels:

```
1 nhanes.samp <- nhanes.samp %>%
2   mutate(
3     HealthGen = factor(HealthGen,
4       levels = c("Excellent",
5                 "Vgood",
6                 "Good",
7                 "Fair",
8                 "Poor")))
9
10 levels(nhanes.samp$HealthGen)
```

```
[1] "Excellent" "Vgood"    "Good"      "Fair"      "Poor"
```



Tip

The `levels` argument defines the order. Categories not listed are set to `NA`.



Cleaning Up Labels with base R

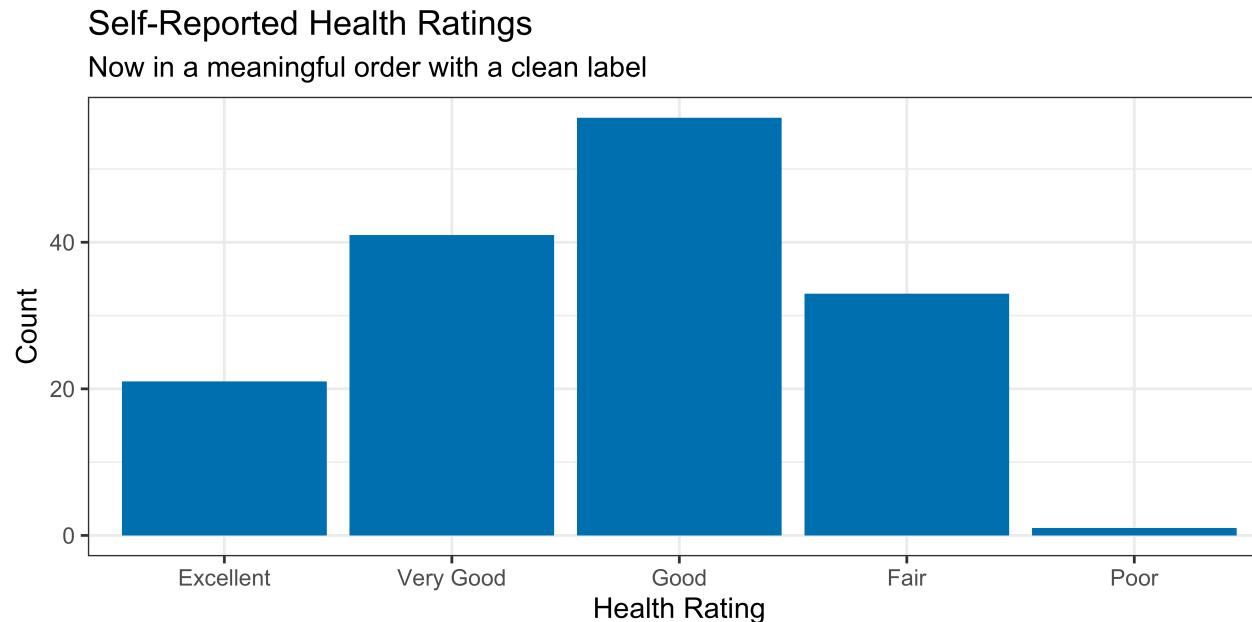
"Vgood" isn't great for a plot or a report. We can rename it:

```
1 # Cleaning Up Labels
2 nhanes.samp <- nhanes.samp %>%
3   mutate(
4     HealthGen = factor(HealthGen,
5       levels = c("Excellent", "Vgood", "Good", "Fair", "Poor"),
6       labels = c("Excellent", "Very Good", "Good", "Fair", "Poor"))
7   )
8
9 levels(nhanes.samp$HealthGen)
[1] "Excellent" "Very Good" "Good"      "Fair"       "Poor"
```



The Fix: Ordered Factor in a Plot

```
1 nhanes.samp %>%
2   filter(!is.na(HealthGen)) %>%
3   ggplot(aes(x = HealthGen)) +
4   geom_bar(fill = "#0072B2") +
5   labs(title = "Self-Reported Health Ratings",
6        subtitle = "Now in a meaningful order with a clean label",
7        x = "Health Rating",
8        y = "Count") +
9   theme_bw(base_size = 16)
```



Same code. The only thing that changed was making `HealthGen` a factor with meaningful levels.



levels() — Inspect and Verify

`levels()` is your diagnostic tool — use it to check that your factor is set up correctly:

```
1 # What levels exist?  
2 levels(nhanes.samp$HealthGen)  
[1] "Excellent" "Very Good" "Good"      "Fair"       "Poor"  
  
1 # How many levels?  
2 nlevels(nhanes.samp$HealthGen)  
[1] 5
```

⚠ Important

Always check `levels()` after creating a factor. If the order isn't what you expect, fix it now — not after you've built plots and run models.



Part 2: Theforcats Package



Why `forcats`?

`factor()` works, but it requires you to type out every level by hand. The `forcats` package (part of the tidyverse) gives you shortcut functions for common tasks:

Function	What it does
<code>fct_relevel()</code>	Manually reorder specific levels
<code>fct_reorder()</code>	Reorder levels by a summary statistic
<code>fct_infreq()</code>	Order levels by frequency
<code>fct_rev()</code>	Reverse the current order
<code>fct_recode()</code>	Rename levels
<code>fct_collapse()</code>	Combine levels into groups

We'll focus on `fct_relevel()` and `fct_reorder()` — the two you'll use most.



Education Is Already a Factor

Education was already a factor in the dataset:

```
1 class(nhanes.samp$Education)
[1] "factor"
1 levels(nhanes.samp$Education)
[1] "8th Grade"      "9 - 11th Grade" "High School"    "Some College"
[5] "College Grad"
```

The levels are in a reasonable order — but what if we want to change the reference category for a model, or reorder by a statistic for a plot?

That's where [forcats](#) comes in.



fct_relevel() — Manual Reordering

Move specific levels to the front, or place them at a specific position:

```
1 # Move "College Grad" to the first position
2 nhanes.samp %>%
3   mutate(Education = fct_relevel(Education, "College Grad")) %>%
4   pull(Education) %>%
5   levels()
```

```
[1] "College Grad"    "8th Grade"      "9 - 11th Grade" "High School"
[5] "Some College"
```

```
1 # Move "College Grad" to the last position
2 nhanes.samp %>%
3   mutate(Education = fct_relevel(Education, "College Grad", after = Inf)) %>%
4   pull(Education) %>%
5   levels()
```

```
[1] "8th Grade"      "9 - 11th Grade" "High School"     "Some College"
[5] "College Grad"
```

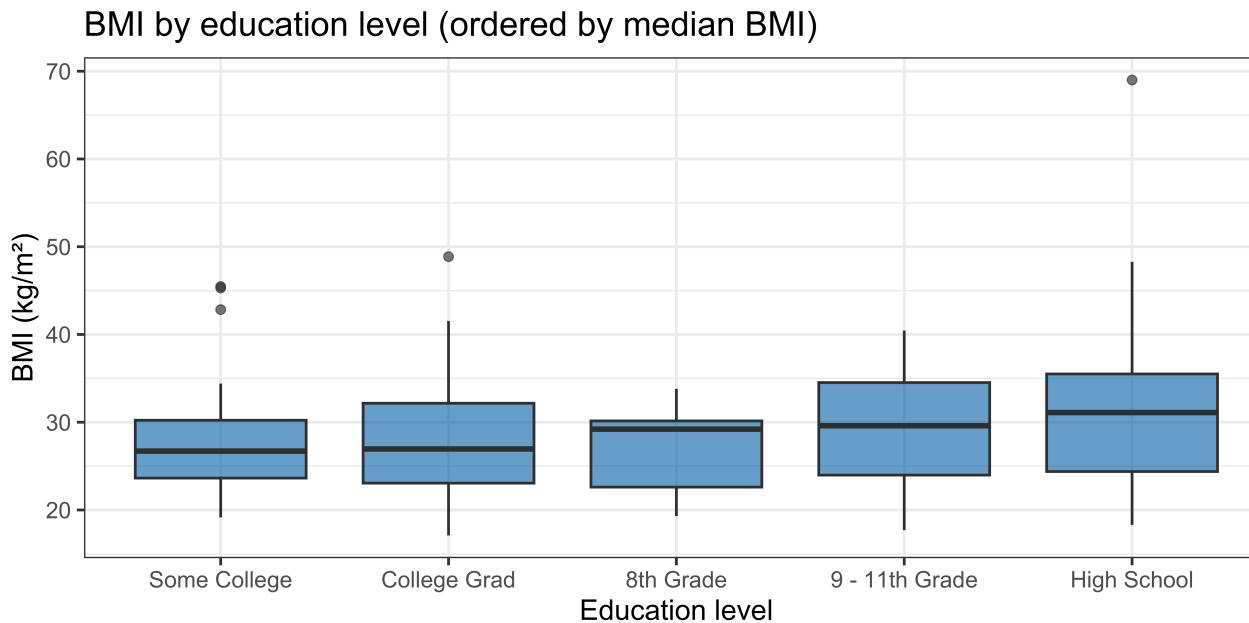
after = Inf is the trick for “put it at the end.”



fct_reorder() – Data-Driven Reordering

Order factor levels by a **summary statistic** of another variable. This is especially powerful for plots:

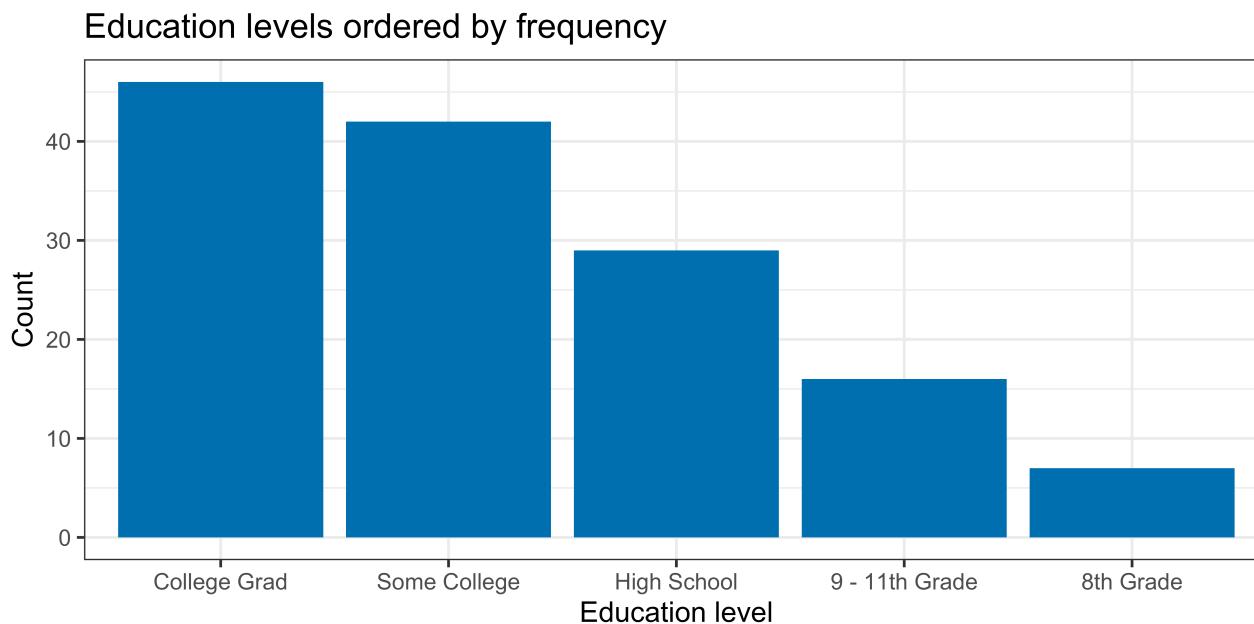
```
1 nhanes.samp %>%
2   filter(!is.na(Education), !is.na(BMI)) %>%
3   mutate(Education = fct_reorder(Education, BMI, .fun = median)) %>%
4   ggplot(aes(x = Education, y = BMI)) +
5   geom_boxplot(fill = "#0072B2", alpha = 0.7) +
6   labs(title = "BMI by education level (ordered by median BMI)",
7        x = "Education level",
8        y = "BMI (kg/m2)") +
9   theme_bw(base_size = 16)
```



fct_infreq() – Order by Frequency

Sometimes you just want categories ordered by how often they appear:

```
1 nhanes.samp %>%
2   filter(!is.na(Education)) %>%
3   mutate(Education = fct_infreq(Education)) %>%
4   ggplot(aes(x = Education)) +
5   geom_bar(fill = "#0072B2") +
6   labs(title = "Education levels ordered by frequency",
7        x = "Education level",
8        y = "Count") +
9   theme_bw(base_size = 16)
```



fct_recode() — Rename Levels

Sometimes level names are too long, abbreviated, or inconsistent. `fct_recode()` lets you rename them:

```
1 unique(nhanes.samp$HealthGen)
[1] Fair      Very Good <NA>      Good      Excellent Poor
Levels: Excellent Very Good Good Fair Poor
1 nhanes.samp %>%
2   mutate(
3     Education = fct_recode(Education,
4                               "Some HS" = "9 - 11th Grade",
5                               "HS Grad" = "High School",
6                               "College" = "College Grad")
7   ) %>%
8   pull(Education) %>%
9   levels()
[1] "8th Grade"    "Some HS"       "HS Grad"      "Some College" "College"
```

Note

- The syntax is `"new name" = "old name"`. It feels backwards at first, but it follows the same `left = right` assignment pattern as `mutate()`. Only the levels you list are renamed — the rest stay as they are.
- Unlike the `labels` argument we used earlier, `fct_recode()` lets you pick and choose. If you only want to rename one category, the others will stay exactly as they were.



Why Factors Matter for Modeling (Preview)

When you fit a regression model, R automatically creates **dummy variables** from factors. The **first level** becomes the **reference category**.

```
1 # Current reference level (first in the list)
2 levels(nhanes.samp$Education)[1]
[1] "8th Grade"
```

```
1 # Change reference to "High School"
2 nhanes.samp %>%
3   mutate(Education = fct_relevel(Education, "High School")) %>%
4   pull(Education) %>%
5   levels()
[1] "High School"     "8th Grade"      "9 - 11th Grade" "Some College"
[5] "College Grad"
```

Note

We'll come back to this in detail when we cover regression. For now, just know that the **first factor level = reference category** in any model.



Factor Tools Summary

Creating factors

- `factor(x, levels = c(...))` — full control
- `levels()` — inspect the result
- Always verify your levels are correct

Common tasks

- Reorder for plots → `fct_reorder()` or `fct_infreq()`
- Set reference level for modeling → `fct_relevel()`
- Rename levels → `fct_recode()`



Tip

The rule of thumb:

- **For plots:** use `fct_reorder()` to order by a statistic, or `fct_infreq()` to order by count
- **For models:** use `fct_relevel()` to set your reference category



Part 3: Wide vs. Long Data



The Conceptual Shift

Most data you receive is in **wide** format — one row per subject, multiple measurement columns.

But many R operations (especially `ggplot2`) need data in **long** format — one row per observation.

Note

This is not about right vs. wrong. It's about which shape your tool expects.



What Does “Wide” Data Look Like?

Each subject is **one row**, with measurements spread across columns:

```
1 nhanes_wide <- nhanes.samp %>%
2   select(ID, Age, BMI, BPSysAve, BPDiaAve) %>%
3   head(5)
4
5 nhanes_wide

# A tibble: 5 × 5
  ID    Age    BMI BPSysAve BPDiaAve
  <int> <int> <dbl>     <int>     <int>
1 63147    41    30        112       76
2 67852    15    19.2       107       72
3 57165    48    27.4       123       69
4 68691     7    17.4        NA       NA
5 69465    50    26.9       152      103
```

Four measurement columns: [Age](#), [BMI](#), [BPSysAve](#), [BPDiaAve](#). This is the shape you’re used to seeing.



What Does “Long” Data Look Like?

Each measurement is **its own row**:

```
1 nhanes_long <- nhanes_wide %>%
2   pivot_longer(
3     cols = c(Age, BMI, BPSysAve, BPDiaAve),
4     names_to = "measure",
5     values_to = "value"
6   )
7
8 nhanes_long
```

```
# A tibble: 20 × 3
  ID measure  value
  <int> <chr>    <dbl>
1 63147 Age      41
2 63147 BMI      30
3 63147 BPSysAve 112
4 63147 BPDiaAve 76
5 67852 Age      15
6 67852 BMI     19.2
7 67852 BPSysAve 107
8 67852 BPDiaAve 72
9 57165 Age      48
10 57165 BMI     27.4
11 57165 BPSysAve 123
12 57165 BPDiaAve 69
13 68691 Age      7
14 68691 BMI     17.4
15 68691 BPSysAve NA
16 68691 BPDiaAve NA
```



Visualizing the Transformation (1/2)

wide

id	x	y	z
1	a	c	e
2	b	d	f

long

id	key	val
1	x	a
2	x	b
1	y	c
2	y	d
1	z	e
2	z	f



Visualizing the Transformation (2/2)

wide

id	x	y	z
1	a	c	e
2	b	d	f



Why Does the Shape Matter?

Wide format is good for:

- Human readability
- Summary tables
- Spreadsheet-style work
- Most statistical tests in R

Long format is good for:

- `ggplot2` — faceting, color mapping
- Grouped operations with `group_by()`
- Repeated measures analysis
- Any time you want to treat variable names as data

Important

ggplot2's golden rule: if you want different variables to be distinguished by color, facet, or grouping, they usually need to be in the **same column** — which means long format.



Part 4: Pivoting



pivot_longer() — Wide to Long

pivot_longer() takes multiple columns and stacks them into two new columns: one for the **variable names** and one for the **values**.

```
1 nhanes.samp %>%
2   select(ID, BMI, BPSysAve, BPDiaAve) %>%
3   head(4) %>%
4   pivot_longer(
5     cols = c(BMI, BPSysAve, BPDiaAve),           # columns to stack
6     names_to = "measurement",                     # new column for variable names
7     values_to = "value"                          # new column for the values
8   )
```

```
# A tibble: 12 × 3
  ID measurement value
  <int> <chr>      <dbl>
1 63147 BMI          30
2 63147 BPSysAve    112
3 63147 BPDiaAve    76
4 67852 BMI         19.2
5 67852 BPSysAve    107
6 67852 BPDiaAve    72
7 57165 BMI         27.4
8 57165 BPSysAve    123
9 57165 BPDiaAve    69
10 68691 BMI        17.4
11 68691 BPSysAve    NA
12 68691 BPDiaAve    NA
```



`pivot_longer()` – The Three Arguments

```
1 pivot_longer(  
2   cols = c(BMI, BPSysAve, BPDiaAve),      # 1. Which columns to pivot  
3   names_to = "measurement",                 # 2. Name for the new "names" column  
4   values_to = "value"                      # 3. Name for the new "values" column  
5 )
```

Argument	Question it answers
<code>cols</code>	Which columns are being stacked?
<code>names_to</code>	What should we call the column that holds the old column names?
<code>values_to</code>	What should we call the column that holds the values?



Think of it as: "I want to make this data **longer** by stacking these columns."



pivot_longer() — Why the Dimensions Change

```
1 nhanes_subset <- nhanes.samp %>%
2   select(ID, BMI, BPSysAve, BPDiaAve)
3
4 # Wide: one row per subject
5 dim(nhanes_subset)
```

```
[1] 200 4
```

```
1 # Long: one row per subject × measurement
2 nhanes_subset %>%
3   pivot_longer(cols = c(BMI, BPSysAve, BPDiaAve),
4                 names_to = "measurement",
5                 values_to = "value") %>%
6   dim()
```

```
[1] 600 3
```

We pivoted **3 columns**, so we get **3x the rows** and go from 4 columns down to 3 (ID + the two new columns).



pivot_wider() — Long to Wide

pivot_wider() does the reverse — it spreads rows back into columns:

```
1 nhanes_long_example <- nhanes.samp %>%
2   select(ID, BMI, BPSysAve, BPDiaAve) %>%
3   head(4) %>%
4   pivot_longer(cols = c(BMI, BPSysAve, BPDiaAve),
5                 names_to = "measurement",
6                 values_to = "value")
7
8 # Now spread it back to wide
9 nhanes_long_example %>%
10  pivot_wider(
11    names_from = "measurement",      # column whose values become new column names
12    values_from = "value"          # column whose values fill the new columns
13  )
# A tibble: 4 × 4
#>   ID    BMI  BPSysAve  BPDiaAve
#>   <int> <dbl>     <dbl>     <dbl>
#> 1 63147   30       112       76
#> 2 67852  19.2      107       72
#> 3 57165  27.4      123       69
#> 4 68691  17.4       NA        NA
```



`pivot_wider()` – The Two Arguments

```
1 pivot_wider(  
2   names_from = "measurement",      # 1. Which column has the new column names?  
3   values_from = "value"           # 2. Which column has the values to fill in?  
4 )
```

Argument	Question it answers
<code>names_from</code>	Which column's values should become new column names?
<code>values_from</code>	Which column's values should fill those new columns?



Tip

Think of it as: "I want to make this data **wider** by spreading this column into multiple columns."



Putting It Together: Pivot → Plot



A Three-Step Workflow

Suppose we want to compare the distributions of `BMI`, `BPSysAve`, and `BPDiaAve` side by side. In wide format, this is awkward. In long format, it's one `ggplot()` call:

Step 1: Select the relevant columns

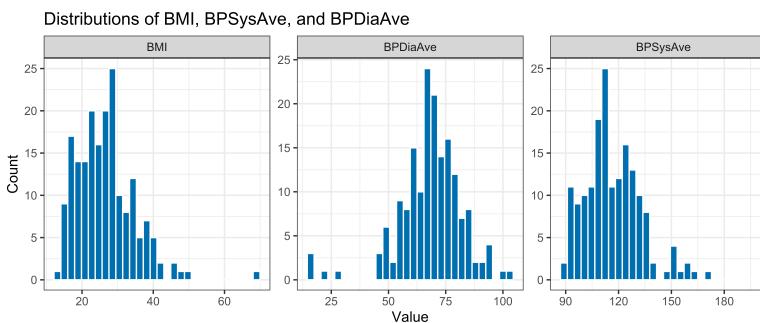
Step 2: Pivot to long format

Step 3: Plot with faceting



Pivot → Plot in Action

```
1 nhanes.samp %>%
2   select(ID, BMI, BPSysAve, BPDiaAve) %>%
3   pivot_longer(
4     cols = c(BMI, BPSysAve, BPDiaAve),
5     names_to = "measurement",
6     values_to = "value"
7   ) %>%
8   ggplot(aes(x = value)) +
9   geom_histogram(bins = 30, fill = "#0072B2", color = "white") +
10  facet_wrap(~ measurement, scales = "free") +
11  labs(title = "Distributions of BMI, BPSysAve, and BPDiaAve",
12        x = "Value",
13        y = "Count") +
14  theme_bw(base_size = 16)
```



One pipeline. Three plots. This is why reshaping matters.



Key Takeaways

Factors

- Characters have no order — factors do
- `factor(x, levels = ...)` for full control
- `fct_reorder()` for data-driven plot ordering
- `fct_relevel()` to set reference categories

Reshaping

- Wide = one row per subject, many measurement columns
- Long = one row per observation
- `pivot_longer()` — wide to long (for plotting, grouping)
- `pivot_wider()` — long to wide (for tables, summaries)



Tip

The rules:

Factors: If your plot categories are in the wrong order, it's a factor problem.

Reshaping: If `ggplot` wants your data in a different shape, pivot it.



References

Factors

- [R for Data Science: Factors](#)
- [forcats documentation](#)
- [Be the boss of your factors](#) — STAT 545

Reshaping data

- [R for Data Science: Pivoting](#)
- [tidyr documentation](#)
- [Garrick Aden-Buie's site](#) — excellent visual explanations

